# Parallel Programming with OpenMP

Jarmo Rantakokko
Senior lecturer, IT UU



```
#pragma      omp parallel
{
  ...
}
```

---
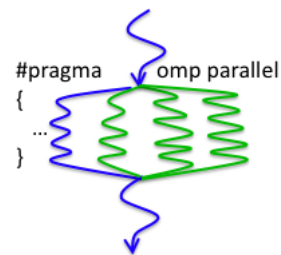
**OpenMP:** Open specification for Multi Processing
(www.openmp.org, v1.0 1997 – v5.0 Nov 8 2018)
Set of compiler directives and functions in C/C++/Fortran

Shared address space model, based on *threads*

**Thread**: - Light weight process, global addresses
- Private program counter, independent
- Private stack pointer, private data

⇒ All threads have access to global data, can run
in parallel and have some private data on stack.

On a multi-core node the threads are scheduled over
the CPU's to the different cores.
⇒ IT-servers Gullviva, Vitsippa, Tussilago have 16
cores and can run 16 parallel threads.

Insert compiler directives for parallelization of
Computations ⇒ high-level model

```
#pragma omp parallel for
for (i=2;i<=N-1;i++)
    A[i]=F(B[i-1]+B[i]+B[i+1]);
```

Loop is automatically parallelized over all
threads, different iterations on different theads.
Arrays A and B are global data, loop variable i
is private.

```
NLOC=N/NPROC
ALLOCATE (A(NLOC),B(NLOC))
. . .
(Standard send/recv avoiding deadlock)

IF (MOD(PID,2)==1) THEN
    CALL MPI_SEND(B(1),LEFT)
    CALL MPI_RECV(TEMP1,LEFT)
ELSEIF (MOD(PID,2)==0 AND PID<NPROC-1)
    CALL MPI_RECV(TEMP2,RIGHT)
    CALL MPI_SEND(B(NLOC),RIGHT)
END IF
IF (MOD(PID,2)==1 AND PID<NPROC-1) THEN
    CALL MPI_SEND(B(NLOC),RIGHT)
    CALL MPI_RECV(TEMP2,RIGHT)
ELSEIF (MOD(PID,2)==0 AND PID>0)
    CALL MPI_RECV(TEMP1,LEFT)
    CALL MPI_SEND(B(1),LEFT)
END IF
                (Simpler with non-blocking communication
                 MPI_Irecv followed by MPI_Send)
IF (PID>0) THEN
    A(1)=F(TEMP1+B(1)+B(2))
END IF
FOR (I=2; I<NLOC-1; I++)
    A(I)=F(B(I-1)+B(I)+B(I+1))
END DO
IF (PID<NPROC-1) THEN
    A(NLOC)=F(B(NLOC-1)+B(NLOC)+TEMP2)
END IF
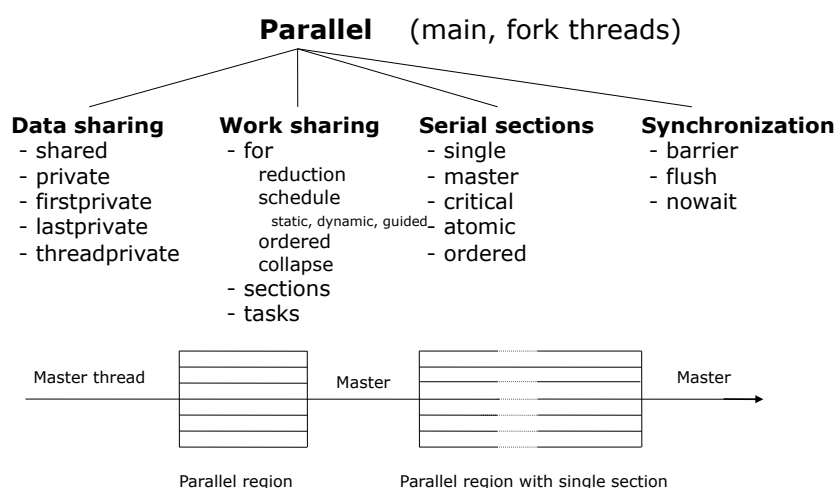```

MPI

```
 struct thread_data
{
        int j1;
        int j2;
};
void *compute(void *arg)
{
        int j,j1,j2;
        struct thread_data *index;
        index=(struct thread_data *)arg;
        j1=index->j1;
        j2=index->j2;

        for (j=j1;j<j2;j++)
                A[j]=F(B[j-1]+B[j]+B[j+1]);
}
int main(){
…
for(t=0; t<NUM_THREADS; t++) {
        index[t].j1=t*len/NUM_THREADS;
        index[t].j2=(t+1)*len/NUM_THREADS;
        pthread_create(&threads[t], &attr, compute,
                        (void *) &index[t]); }
```

Pthreads

# OpenMP directives:

**Parallel**    (main, fork threads)

**Data sharing**
- shared
- private
- firstprivate
- lastprivate
- threadprivate

**Work sharing**
- for
    reduction
    schedule
        static, dynamic, guided
    ordered
    collapse
- sections
- tasks

**Serial sections**
- single
- master
- critical
- atomic
- ordered

**Synchronization**
- barrier
- flush
- nowait

Master thread

Master

Master

Parallel region

Parallel region with single section

## OpenMP library functions:
- omp_set_num_threads
- omp_get_num_threads
- omp_get_max_threads
- omp_get_thread_num
- omp_set_nested
- and more (e.g. lock)

Allows for more flexible and user controlled (e.g. load balancing) programming than with the standard directives.

**Environment variables:**   (export VARIABLE=value)
- OMP_NUM_THREADS
- OMP_SCHEDULE
- OMP_NESTED
- and more (stacksize, wait policy)

To run on 4 threads, before start of program do:
export OMP_NUM_THREADS=4

---

## Directives:  (Support only in Fortran/C/C++)

C/C++:    #pragma omp *directive*
          {    code block    }

Fortran:    !$omp *directive*
                code block
            !$omp end *directive*

**Note**: The directives are ignored by non-supporting compiler or if OpenMP-flag is turned off in compiling.
⇒ Portable code between single CPU, multi-core, and general parallel computers.

Also, possible to parallelize code incrementally
(start with heaviest routine and continue until sufficient parallelism and performance are achieved)

**Parallel:** (Fork-Join of threads)

```
#pragma omp parallel [subdirectives]
{
"parallel code"
}
```

Subdirectives:

```
if ( true/false )      --  parallel/serial
num_threads( int )   --  Number of threads
reduction (op:var)    --  parallel reduction
```

+ directives for data sharing (private/shared)

---

**Parallel:** (Fork-Join of threads)

```
#pragma omp parallel
{
"parallel code"
}
```

If no subdirectives, all data shared (global) and all code executed in parallel by all threads. At the end of parallel the threads are synchronized and joined.

```
Ex: program p1
    ...
    call fred()
    #pragma omp parallel
    { call billy() }
    call daisy()
```

master

fred()

billy() – all execute

daisy()
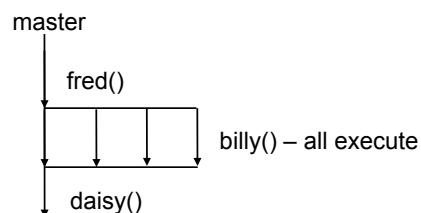
**Example HelloWorld:**

```c
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[]) {

#pragma omp parallel
{
    printf("Hello World! %d\n",omp_get_thread_num());
}

}
```

**Task:** Compile and run the program helloworld.c
➢ gcc –fopenmp helloworld.c –o hello
➢ ./hello
Run on different number of threads. In what ways can
we change the number of threads?What is the default
number of threads? What is the max number of threads?

# Data sharing:

• **shared**( [list of variables] )  - default
• **private**( [list of variables] )

```
Ex: program p2
    ...
    a=100; b=0;
    #pragma omp parallel private(a) num_threads(10)
    {
        b=b+1000;
        a=b+a;
    }
    printf("a= %d, b= %d \n", a,b);
```

**Task:** What is the result, run the program datasharing.c
several times and explain the output.

**Note**: All private variables are allocated on the stack
=> <u>uninitialized</u> at entry and removed at exit,
*original a* not equal to *private a*!


**Note2:** Shared variables must be protected from
simultaneous writes by different threads!
(Use critical directive or locks.)

---

• **firstprivate**( [list of variables] )
As private but the variables are initialized from the
original variable (in master) before parallel.


• **lastprivate**( [list of variables] )
At exit, the original variable gets the value from the
thread executing the last iteration in a <u>loop</u> using the
for-directive or the last <u>section</u> in the sections-directive.


• **threadprivate**( [list of variables] )
Make global file scope variables local and persistent
to a thread through the execution of multiple parallel
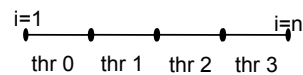regions.

## Work sharing (within parallel)
- Loop level parallelism – **for**
- Task parallelism – **sections, tasks**
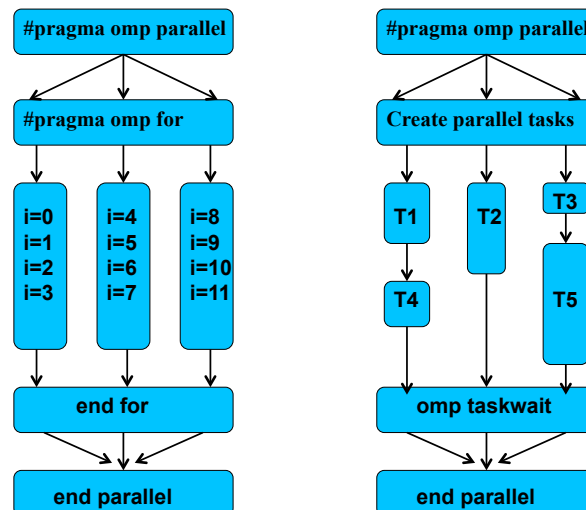
---

## for-directive:

```
#pragma omp for [subdirectives]
for (i=1; i<=n; i++)
  { loop-body }
```



**Subdirectives:**
- Private
- Firstprivate
- Lastprivate

- Reduction
- Schedule
- Ordered
- Collapse

Without subdirectives, loop counter is private, loop space is divided statically into nthr equal pieces, and run in parallel (different iterations in different threads). Threads are synchronized at end of the for-directive.

**Note:** We must have a perfectly parallel loop!

## Example: Enumeration sort

```
for (j=0;j<len;j++)
{
  rank=0;
  for (i=0;i<len;i++)
    if (indata[i]<indata[j]) rank++;
  outdata[rank]=indata[j];
}
```

For each element (j) check how many other
elements (i) are smaller than it => rank
Perfectly parallel tasks for each element (j)

**Task:** Parallelize the j-loop in enumsort.c and set
appropriate variables as private.
What speedup can you get for 50,000 elements?

---

## Reduction( op:[list of variables] )
Performs a global reduction using
**op**=+,-,*,max,min, or a logical operator

```
sum=0;
#pragma omp parallel for reduction(+:sum)
for (i=0;i<n;i++) sum=sum+a[i];
```

```
sum=0;
#pragma omp parallel private(locsum)
{
    locsum=0;
    #pragma omp for
    for (i=0;i<n;i++) locsum=locsum+a[i];
    #pragma omp critical
    { sum=sum+locsum; }
}
```

**Task:** Parallelize the inner i-loop in enumsort.c and
compare the performance with your first parallelization.
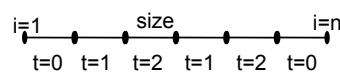What are the performance obstacles and/or advantages?

**Schedule( type, [size] )**
Divides the iteration space into chunks=size and
schedules the chunks to threads according to type.
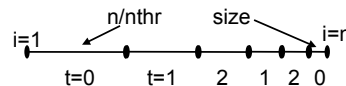(size=n/nthr by default)

**type=static:**
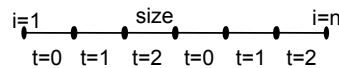


Assign the chunks
cyclicly to threads

**type=dynamic:**



Dynamic scheduling, as
soon as a thread is ready
it gets a new chunk

**type=guided:**



As dynamic but the chunk size
is decreasing towards end.
Minimizes synchronization time.

---

**type=runtime:**

Decide at runtime using the environment variable
export *schedule=type* (where *type* is some above).

**type=auto:**

Let the run-time system and/or compiler decide automatically.

**Note:** Static scheduling is good for data locality (cache)
while dynamic/guided good for load balance.

**Task:** Try different scheduling options in the program
loop.c, what gives the best performance, what is the
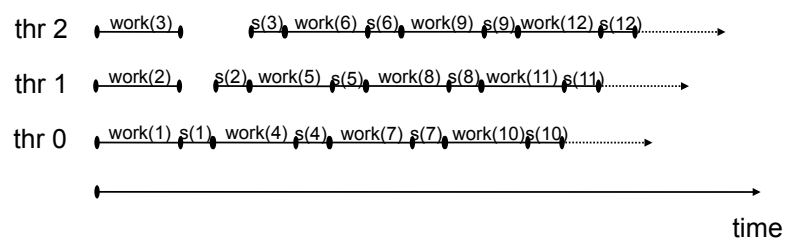theoretically minimal runtime, how can we get that?

## Ordered

Only one thread is allowed to the ordered block at a time and sequentially in loop order. Useful for I/O.

```
#pragma omp parallel
{
    #pragma omp for schedule(static,1) ordered
    for (i=0;i<n;i++)
    {
        call work(i)            ! parallel work

        #pragma omp ordered
        { call s(i) }           ! serial section
    }
}
```

Assume work(i)>>s(i) => Parallelism, pipelining effect

---

## Static,1:

thr 2 • work(3) • — s(3)• work(6) • s(6)• work(9) • s(9)•work(12) • s(12).............→

thr 1 • work(2) • — s(2)• work(5) • s(5)• work(8) • s(8)•work(11) • s(11).................→

thr 0 • work(1) • s(1)• work(4) • s(4)• work(7) • s(7)•work(10)•s(10).......................→

•————————————————————————————→
                                                        time

What happens if we use default scheduling (size=n/nthr)?

## Collapse directive

Allow *collapsing* of perfectly nested loops, i.e.,
form a single loop and then parallelize that.
Example: parallelize both i and j-loop in MxM

```
#pragma omp parallel
{
#pragma omp for collapse(2) private(i,j,k)
for (i=0;i<len;i++)
   for (j=0;j<len;j++)
   {
      c[i*len+j]=0.0;
      for (k=0;k<len;k++)
         c[i*len+j]+=a[i*len+k]*b[k*len+j];
   }
}
```

## Task parallelism (static predefined tasks)

### Sections

```
#pragma omp sections [subdirectives]
{
    #pragma omp section
        { task 1 }

    #pragma omp section
        { task 2 }

    etc.
}
```

Subdirectives:
• Private
• Firstprivate
• Lastprivate
• Reduction

The sections/tasks are scheduled (statically) to the
threads and run in parallel. At end of sections the
threads are synchronized. (No load balancing).

**Nested parallelism** (load balancing of sections)

```
omp_set_nested(1);
#pragma omp parallel sections num_threads(2)
{
    #pragma omp section
    {
    #pragma omp parallel for num_threads(P1)
    for (k=0;k<n1;k++)
        call WORK1(A[K])
    }
    #pragma omp section
    {
    #pragma omp parallel for num_threads(P2)
    for (k=0;k<n2;k++)
        call WORK2(A[K])
    }

}
```

Assign appropriate number of threads to each section.

**Task:** Make a two-level parallelization of enumsort.c

---

# Task directive (dynamic tasks)

Can implement task-queues that are scheduled dynamically to all available threads in a parallel environment. Tasks can be generated at run-time.

Generate a task:
```
#pragma omp task [if/untied/'datasharing']
{ task }
```

Wait for all tasks to complete
```
#pragma omp taskwait
```

Task scheduling points at following locations:
1. Generation of task
2. Last instruction in task
3. Taskwait-directive
4. Implicit and explicit barriers

**Task:** Study and run the program task.c. What is the effect of if and nowait?

## Serial sections

Avoid terminating threads, loose data in cache
if threads rescheduled to different CPUs or cores
(with fork-join model).

**#pragma omp single [subdirectives]**
The code-block within single is executed only by one
thread, the others skip and wait at the end of block.
Subdirectives: - private
                 - firstprivate

**#pragma omp master**
The code-block is executed only by master thread,
the other skip and continue (no barrier).

---

**#pragma omp critical [name]**
The code-block is executed by one thread at a time.
As ordered but no predefined order.
If no name all critical sections have the same name.
Only one critical section with the same name can be
executed by one thread at a time.

**#pragma omp atomic**
Atomic update by one thread at a time. As critical
but applies only for a one line expression. (Does
not include a memory flush.)

## Synchronization

Done implicitly at end of:
- parallel
- for
- sections
- single

Explicit barrier:
#pragma omp barrier

Can override with *nowait:*

    #pragma omp for nowait
    for (i=0;i<n;i++)
        { code }

**Note:** If *nowait* be careful not to use data updated by other threads, *nowait* overrides memory flush!

---

In a **memory flush** all thread visible shared variables are refreshed (caches invalidated and memory updated). A memory flush is performed at all barriers (implicit and explicit) and before/after a critical directive.



Conceptual model
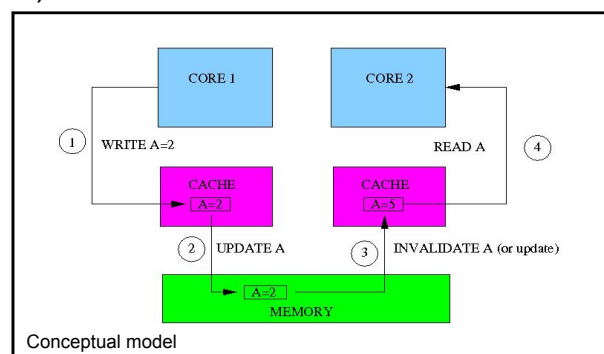
=> *Be very careful with nowait !!! (Nowait removes 2 & 3)*

OpenMP has a *relaxed-consistency* model, i.e., the threads can cache data not keeping exact consistency.

**Task:** Run the program memory.c and explain
its output. How can we fix the problem?

```c
int main(int argc, char *argv[]) {
int id,nthr;

#pragma omp parallel private(id)
  {
    nthr=-1;
    id=omp_get_thread_num();

    #pragma omp single
    { nthr=omp_get_num_threads(); }

    printf("Hello World! %d %d\n",id,nthr);
  }
}
```

## Synchronization with locks

Can lock a code section and/or data only accessible
to a specific thread. Routines include a flush.

```
omp_init_lock - omp_destroy_lock
omp_set_lock  - omp_unset_lock
omp_test_lock
```

Example:
```
omp_lock_t lockvar;
omp_init_lock(lockvar);
…
#pragma omp parallel {
…
    omp_set_lock(lockvar);
    sum=sum+a;
    omp_unset_lock(lockvar);
}
```